


CHL7001H S1 Applied Deep Learning


Modern Deep Learning

Recap of last week

- Fully connected layers (e.g. Multilayer Perceptrons)
- In the simplest case, all input units are connected to all output units. We call this a **fully connected layer**.
- A multilayer network consisting of fully connected layers is called a **multilayer perceptron**.




Some activation functions




Identity

$$f(x) = x$$




ReLU

$$f(x) = \max(0, x)$$



Logistic

$$f(x) = \frac{1}{1 + e^{-x}}$$



TanH

$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



Example of shallow NN

- <https://playground.tensorflow.org/#activation=tanh&batchSize=10&dataset=spiral®Dataset=reg-plane&learningRate=0.03®ularizationRate=0&noise=0&networkShape=8&seed=0.50265&showTestData=false&discretize=false&percTrainData=50&x=true&y=true&xTimesY=false&xSquared=false&ySquared=false&cosX=false&sinX=false&cosY=false&sinY=false&collectStats=false&problem=classification&initZero=false&hideText=false>

Example with a deep NN

- <https://playground.tensorflow.org/#activation=relu&batchSize=30&dataset=spiral®Dataset=reg-plane&learningRate=0.1®ularizationRate=0&noise=0&networkShape=4,4,4&seed=0.50265&showTestData=false&discretize=false&percentTrainData=50&x=true&y=true&xTimesY=false&xSquared=false&ySquared=false&cosX=false&sinX=false&cosY=false&sinY=false&collectStats=false&problem=classification&initZero=false&hideText=false>

Handwritten digits classification




Problem with fully connected network


What is the problem with having this as the architecture (densely connected)?

- Too many parameters!
 - Input: $28*28 = 784$
 - If we plan to have 100 hidden units, then the number of parameters will be $784*100 = 78.4K$
- What happens if the object in the image shifts a little?
- No spatial information!

Example of convolution




Using spatial structure



- Connect patch in input layer to a single neuron in subsequent layer.
- Use a sliding window to define connections.

Feature extraction with convolution



- Filter of size 4x4 : 16 different weights
- Apply this same filter to 4x4 patches in input
- Shift by 2 pixels for next patch

This “patchy” operation is convolution

1. Apply a set of weights – a filter – to extract local features
2. Use multiple filters to extract different features - sliding
3. Spatially share parameters of each filter or say weights are shared between all image location

Feature maps



Original



Sharpen




Edge Detect



"Strong" Edge
Detect

Convolution operation



Convolution operation

Suppose we want to compute the convolution of a 5x5 image and a 3x3 filter:

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

image




1	0	1
0	1	0
1	0	1


filter

Convolution operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:



Convolution operation




Simple filters

0	0	0	0	0
0	$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$	0
0	$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$	0
0	$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$	0
0	0	0	0	0



Simple filters

0	0	0	0	0
0	0	0	0	0
0	-1	1	0	0
0	0	0	0	0
0	0	0	0	0



Feature maps



Original



Sharpen

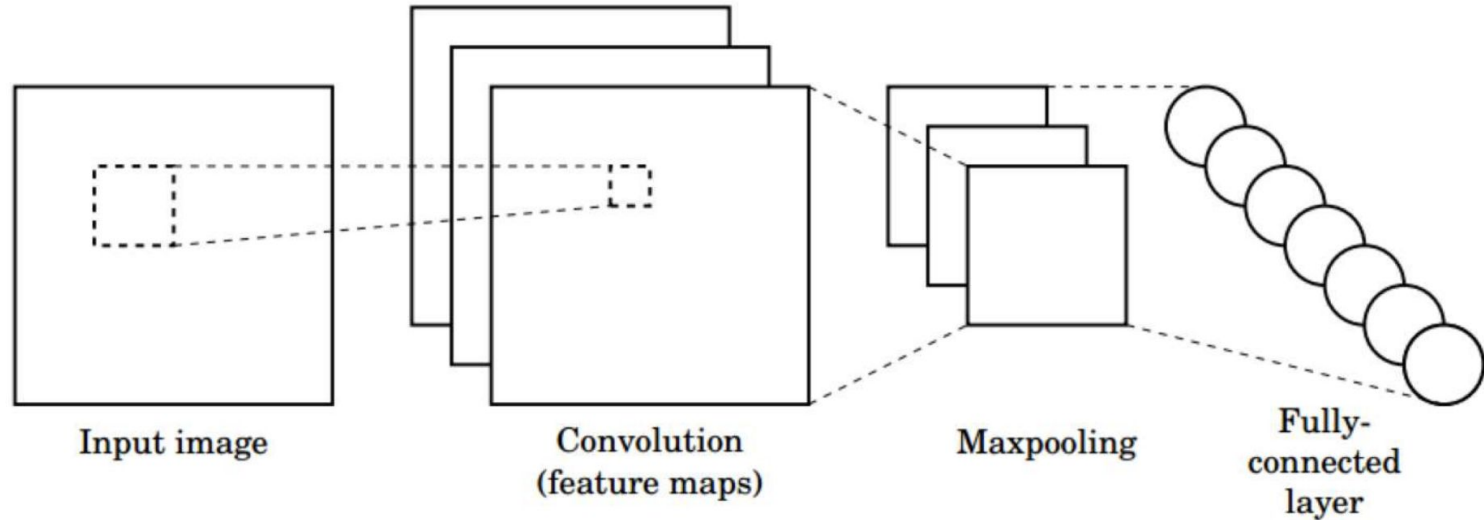


Edge Detect




"Strong" Edge
Detect

CNN for classification



1. Convolution: Apply filters with learned weights to generate feature maps.
2. Non-linearity: Often ReLU.
3. Pooling: Downsampling operation on each feature map.

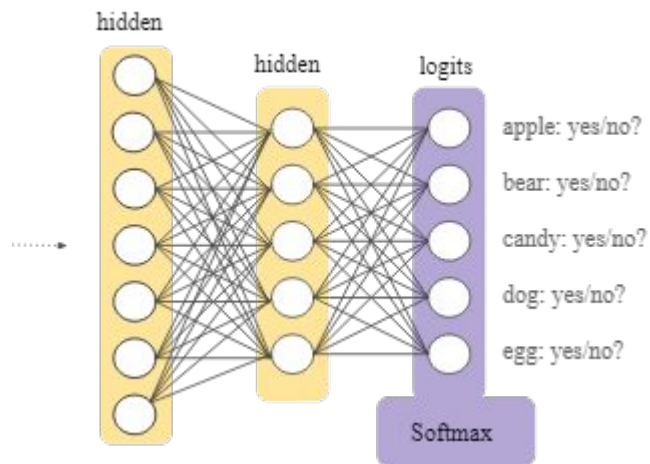
CNN for classification



1. Convolution: Apply filters with learned weights to generate feature maps.
2. Non-linearity: Often ReLU.
3. Pooling: Downsampling operation on each feature map.

Multiclass

- Multiclass, single label? Softmax
- Multiclass, multilabel? Sigmoids



Examples

- <https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/r2/tutorials/quickstart/beginner.ipynb>
- https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/r2/tutorials/images/intro_to_cnns.ipynb